



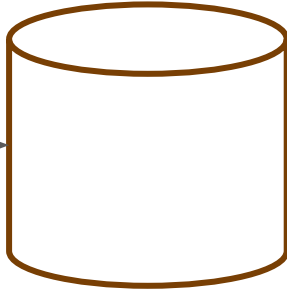
# INFOH417 Database System Architectures

Mahmoud SAKR <[mahmoud.sakr@ulb.be](mailto:mahmoud.sakr@ulb.be)>

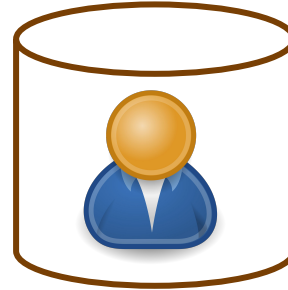
École polytechnique de Bruxelles

2021

# What is this course for ?



SQL  
Python  
Web  
...



Storage  
Access control  
Optimization  
Distribution  
...

# Course Goals

- Understanding the query optimization and execution cycle
- Improving slow queries
- Describing the common index structures, knowing their capabilities and shortcomings
- Understanding cost based optimization, and the associated statistics and estimation methods
- Describing and being able to implement Abstract Data Types in extensible database systems
- Describing data and query distribution mechanisms, and being able to configure and run a distributed database system
- Comparing the architectures of SQL and NoSQL systems
- Comparing the architectures of database V.S. data flow systems

# Course Topics

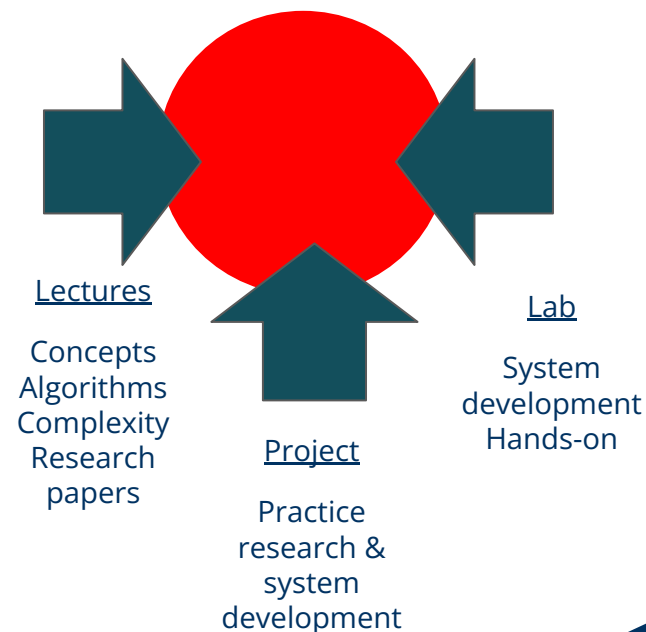
- Query Execution
- Refreshing SQL and Relational Algebra
- Query Optimization
- Indexes
- Advanced Databases
- Multidimensional index structures
- Distributed databases
- New DBMS Architectures (NoSQL)
- New Data Flow Architectures (Map-Reduce and Spark)

# Prerequisites

- Relational databases
- SQL
- Relational Algebra
- General programming skills

# Course Organization

- Lectures on Friday, S.UB4.132, 10h00-12h00
- Lab sessions on Thursday, S.UB4.130, 6x[10h00-12h00], 5x[16h00-18h00]
  - Taught by: Maxime Schoemans
  - make PostgreSQL from [source](#)
- Project
  - Implement data management features in PostgreSQL
- Check [timeedit](#) for schedule updates
- Grading
  - Group project, 4 members, 40%
  - Written exam, 60%
- Course notes, please enroll in [Université virtuelle](#)



# Recommended Readings

- A mixture of book chapters and research papers, which will be identified per lecture

# Refreshing SQL

- Silberschatz, Korth and Sudarshan, Database Systems Concepts 6th Edition, Chapter 3.



# Course Schema

| instructor            |
|-----------------------|
| <u>ID</u> char(5)     |
| name varchar(20)      |
| dept_name varchar(20) |
| salary numeric(8, 2)  |

| takes                       |
|-----------------------------|
| <u>ID</u> char(5)           |
| <u>courseID</u> varchar(20) |
| <u>semester</u> varchar(6)  |
| <u>year</u> numeric(4,0)    |
| grade varchar(2)            |

| student               |
|-----------------------|
| <u>ID</u> char(5)     |
| name varchar(20)      |
| dept_name varchar(20) |
| tot_cred numeric(3,0) |

| course                  |
|-------------------------|
| <u>courseID</u> char(5) |
| title varchar(50)       |
| dept_name varchar(20)   |
| credits numeric(2,0)    |

```
create table takes (  
    ID                varchar(5),  
    course_id         varchar(8),  
    sec_id            varchar(8),  
    semester          varchar(6),  
    year              numeric(4,0),  
    grade             varchar(2),  
    primary key (ID, course_id, semester, year) ,  
    foreign key (ID) references student,  
    foreign key (course_id) references course);
```

# Updates to tables

- **Insert**

`insert into instructor values ('10211', 'Smith', 'Biology', 66000);`

- **Delete**

- Remove all tuples from the *student* relation

`delete from student`

- **Drop Table**

`drop table r`

- **Alter**

`alter table r add A D`

- where *A* is the name of the attribute to be added to relation *r* and *D* is the domain of *A*.
- All existing tuples in the relation are assigned *null* as the value for the new attribute.

`alter table r drop A`

- where *A* is the name of an attribute of relation *r*
- Dropping of attributes not supported by many databases.

# The select Clause

Find the department names of all instructors, and remove duplicates

```
select distinct dept_name  
from instructor
```

The keyword **all** specifies that duplicates should not be removed.

```
select all dept_name  
from instructor
```

## The select Clause (Cont.)

An asterisk in the select clause denotes “all attributes”

```
select *  
from instructor
```

An attribute can be a literal with no **from** clause

```
select '437'  
select '437' as F00
```



## The select Clause (Cont.)

The **select** clause can contain arithmetic expressions involving the operation, +, −, \*, and /, and operating on constants or attributes of tuples.

```
select ID, name, salary/12 as monthlySalary  
from instructor
```

# The where Clause

To find all instructors in Comp. Sci. dept

```
select name  
from instructor  
where dept_name = 'Comp. Sci.'
```

Logical connectives **and, or, and not**

To find all instructors in Comp. Sci. dept with salary > 80000

# The where Clause

To find all instructors in Comp. Sci. dept

```
select name  
from instructor  
where dept_name = 'Comp. Sci.'
```

Logical connectives **and**, **or**, and **not**

To find all instructors in Comp. Sci. dept with salary > 80000

```
select name  
from instructor  
where dept_name = 'Comp. Sci.' and salary > 80000
```

## The from Clause

Find the names of all instructors in the Art department who have taught some course and the course\_id

```
select name, course_id
from instructor , teaches
where instructor.ID = teaches.ID and instructor.
dept_name = 'Art'
```



# Self Join Example

Relation *emp-super*

| person | supervisor |
|--------|------------|
| Bob    | Alice      |
| Mary   | Susan      |
| Alice  | David      |
| David  | Marry      |

Find the supervisor of “Bob”

Find the supervisor of the supervisor of “Bob”

# String Pattern matching

SIMILAR TO is similar to LIKE, and further supports regex operators

```
SELECT 'abc' LIKE 'abc';      true
```

```
SELECT 'abc' LIKE 'a%';      true
```

```
SELECT 'abc' LIKE '_b_';      true
```

```
SELECT 'abc' LIKE 'c';       false
```

```
SELECT 'abc' SIMILAR TO 'abc';      true
```

```
SELECT 'abc' SIMILAR TO 'a';        false
```

```
SELECT 'abc' SIMILAR TO '%(b|d)%';  true
```

```
SELECT 'aaac' SIMILAR TO 'a*b*c%';   true
```

```
SELECT 'aaac' SIMILAR TO 'a*b+c%';   false
```

<https://www.postgresql.org/docs/13/functions-matching.html>

# Ordering the Display of Tuples

List in alphabetic order the names of all instructors

```
select distinct name  
from   instructor  
order by name
```

```
order by name desc
```

```
order by dept_name, name
```

# Where Clause Predicates

Example: Find the names of all instructors with salary between \$90,000 and \$100,000 (that is,  $\geq \$90,000$  and  $\leq \$100,000$ )

```
select name
from instructor
where salary between 90000 and 100000
```

Tuple comparison

```
select name, course_id
from instructor, teaches
where (instructor.ID, dept_name) = (teaches.ID, 'Biology');
```

# Set Operations

Find courses that ran in Fall 2009 or in Spring 2010

```
(select course_id from section where sem = 'Fall' and year = 2009)
union      ||
(select course_id from section where sem = 'Spring' and year = 2010)
```

Find courses that ran in Fall 2009 and in Spring 2010

```
(select course_id from section where sem = 'Fall' and year = 2009)
intersect
(select course_id from section where sem = 'Spring' and year = 2010)
```

Find courses that ran in Fall 2009 but not in Spring 2010

```
(select course_id from section where sem = 'Fall' and year = 2009)
except
(select course_id from section where sem = 'Spring' and year = 2010)
```

## Set Operations (Cont.)

Set operations **union**, **intersect**, and **except**

Each of the above operations automatically eliminates duplicates

To retain all duplicates use the corresponding multiset versions **union all**, **intersect all** and **except all**.

Suppose a tuple occurs  $m$  times in  $r$  and  $n$  times in  $s$ , then, it occurs:

? times in  $r$  **union all**  $s$

? times in  $r$  **intersect all**  $s$

? times in  $r$  **except all**  $s$

## Set Operations (Cont.)

Set operations **union**, **intersect**, and **except**

Each of the above operations automatically eliminates duplicates

To retain all duplicates use the corresponding multiset versions **union all**, **intersect all** and **except all**.

Suppose a tuple occurs  $m$  times in  $r$  and  $n$  times in  $s$ , then, it occurs:

$m + n$  times in  $r$  **union all**  $s$

$\min(m, n)$  times in  $r$  **intersect all**  $s$

$\max(0, m - n)$  times in  $r$  **except all**  $s$

# Null Values

*null* signifies an unknown value or that a value does not exist.

The result of any **arithmetic** expression involving *null* is *null*

Example: **5 + null** returns null

The predicate **is null** can be used to check for null values.

Example: Find all instructors whose salary is null.

```
select name  
from instructor  
where salary is null
```



# Null Values and Three Valued Logic

| NULL  | AND   | OR                    | NOT  |
|-------|-------|-----------------------|------|
| TRUE  | NULL  | TRUE                  | NULL |
| FALSE | FALSE | <del>FALSE</del> Null |      |
| NULL  | NULL  | NULL                  |      |

# Aggregate Functions

These functions operate on the multiset of values of a column of a relation, and return a value

**avg:** average value

**min:** minimum value

**max:** maximum value

**sum:** sum of values

**count:** number of values

# Aggregate Functions – Group By

Find the average salary of instructors in each department

```
select dept_name, avg (salary) as avg_salary  
from instructor  
group by dept_name;
```

| ID    | name       | dept_name  | salary |
|-------|------------|------------|--------|
| 76766 | Crick      | Biology    | 72000  |
| 45565 | Katz       | Comp. Sci. | 75000  |
| 10101 | Srinivasan | Comp. Sci. | 65000  |
| 83821 | Brandt     | Comp. Sci. | 92000  |
| 98345 | Kim        | Elec. Eng. | 80000  |
| 12121 | Wu         | Finance    | 90000  |
| 76543 | Singh      | Finance    | 80000  |
| 32343 | El Said    | History    | 60000  |
| 58583 | Califieri  | History    | 62000  |
| 15151 | Mozart     | Music      | 40000  |
| 33456 | Gold       | Physics    | 87000  |
| 22222 | Einstein   | Physics    | 95000  |

| dept_name  | avg_salary |
|------------|------------|
| Biology    | 72000      |
| Comp. Sci. | 77333      |
| Elec. Eng. | 80000      |
| Finance    | 85000      |
| History    | 61000      |
| Music      | 40000      |
| Physics    | 91000      |

## Aggregation (Cont.)

Attributes in **select** clause outside of aggregate functions must appear in **group by** list

```
/* erroneous query */  
select dept_name, ID, avg (salary)  
from instructor  
group by dept_name;
```

Why ?

To have one average for each group formed by group by.

# Aggregate Functions – Having Clause

Find the names and average salaries of all departments whose average salary is greater than 42000

```
select dept_name, avg (salary) as moyen
from instructor
group by dept_name
having avg (salary) > 42000;
```

Note: predicates in the **having** clause are applied after the formation of groups whereas predicates in the **where** clause are applied before forming groups

# Null Values and Aggregates

Total all salaries

```
select sum (salary )  
from instructor
```

Above statement ignores null amounts

Result is *null* if there is no non-null amount

All aggregate operations except **count(\*)** ignore tuples with null values on the aggregated attributes

What if collection has only null values?

count returns 0



all other aggregates return null

# Nested Subqueries

The nesting can be done in the following SQL query

```
select  $A_1, A_2, \dots, A_n$   
from  $r_1, r_2, \dots, r_m$   
where  $P$ 
```

as follows:

$A_i$  can be replaced by a subquery that generates a single value.

$r_i$  can be replaced by any valid subquery

$P$  can be replaced with an expression of the form:

$B <\text{operation}> (\text{subquery})$

Where  $B$  is an attribute and  $<\text{operation}>$  to be defined later.

# Set Membership

Find courses offered in Fall 2009 and in Spring 2010

```
select distinct course_id
from section
where semester = 'Fall' and year= 2009 and
       course_id in (select course_id
                     from section
                     where semester = 'Spring' and year= 2010);
```

Find courses offered in Fall 2009 but not in Spring 2010



# Set Membership

Find courses offered in Fall 2009 and in Spring 2010

```
select distinct course_id
from section
where semester = 'Fall' and year= 2009 and
       course_id in (select course_id
                      from section
                      where semester = 'Spring' and year= 2010);
```

Find courses offered in Fall 2009 but not in Spring 2010

```
select distinct course_id
from section
where semester = 'Fall' and year= 2009 and
       course_id not in (select course_id
                          from section
                          where semester = 'Spring' and year= 2010);
```

## Set Comparison – “some” Clause

Find names of instructors with salary greater than that of some (at least one) instructor in the Biology department.

```
select distinct T.name
from instructor as T, instructor as S
where T.salary > S.salary and S.dept name = 'Biology';
```

Same query using > **some** clause

```
select name
from instructor
where salary > some (select salary
                      from instructor
                      where dept name = 'Biology');
```

# Definition of “some” Clause

- $F <\text{comp}> \text{some } r \Leftrightarrow \exists t \in r \text{ such that } (F <\text{comp}> t)$   
Where  $<\text{comp}>$  can be:  $<, \leq, >, =, \neq$

$(5 < \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline 6 \\ \hline \end{array}) = \text{true}$  (read: 5 < some tuple in the relation)

$(5 < \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{false}$

$(5 = \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{true}$

$(5 \neq \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{true (since } 0 \neq 5)$

$(= \text{some}) \equiv \text{in}$

However,  $(\neq \text{some}) \not\equiv \text{not in}$

## Set Comparison – “all” Clause

Find the names of all instructors whose salary is greater than the salary of all instructors in the Biology department.

## Set Comparison – “all” Clause

Find the names of all instructors whose salary is greater than the salary of all instructors in the Biology department.

```
select name
from instructor
where salary > all (select salary
                    from instructor
                    where dept name = 'Biology');
```

# Definition of “all” Clause

$F \text{ <comp> all } r \Leftrightarrow \forall t \in r (F \text{ <comp> } t)$

$(5 \text{ < all } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline 6 \\ \hline \end{array}) = \text{false}$

$(5 \text{ < all } \begin{array}{|c|} \hline 6 \\ \hline 10 \\ \hline \end{array}) = \text{true}$

$(5 = \text{all } \begin{array}{|c|} \hline 4 \\ \hline 5 \\ \hline \end{array}) = \text{false}$

$(5 \neq \text{all } \begin{array}{|c|} \hline 4 \\ \hline 6 \\ \hline \end{array}) = \text{true (since } 5 \neq 4 \text{ and } 5 \neq 6)$

**$(\neq \text{ all}) \equiv \text{not in}$**

However,  **$(= \text{ all}) \not\equiv \text{in}$**

# Test for Empty Relations

- The **exists** construct returns the value **true** if the argument subquery is nonempty.
- **exists**  $r \Leftrightarrow r \neq \emptyset$
- **not exists**  $r \Leftrightarrow r = \emptyset$

## Use of “exists” Clause

Yet another way of specifying the query “Find all courses taught in both the Fall 2009 semester and in the Spring 2010 semester”

```
select course_id
from section as S
where semester = 'Fall' and year = 2009 and
exists ???
```



## Use of “exists” Clause

Yet another way of specifying the query “Find all courses taught in both the Fall 2009 semester and in the Spring 2010 semester”

```
select course_id
from section as S
where semester = 'Fall' and year = 2009 and
    exists (select *
            from section as T
            where semester = 'Spring' and year= 2010
            and S.course_id = T.course_id);
```

**Correlation name** – variable S in the outer query

**Correlated subquery** – the inner query

## Use of “not exists” Clause

Find all students who have taken all courses offered in the Biology department.

```
select distinct S.ID, S.name
from student as S
where not exists ( (select course_id
                    from course
                    where dept_name = 'Biology')
                  except
                  (select T.course_id
                   from takes as T
                   where S.ID = T.ID));
```

- First nested query lists all courses offered in Biology
- Second nested query lists all courses a particular student took

Note that  $X - Y = \emptyset \Leftrightarrow X \subseteq Y$

*Note:* Cannot write this query using = **all** and its variants

# Test for Absence of Duplicate Tuples

The **unique** construct tests whether a subquery has any duplicate tuples in its result.

The **unique** construct evaluates to “true” if a given subquery contains no duplicates .

Find all courses that were offered at most once in 2009

```
select T.course_id
from course as T
where unique (select R.course_id
              from section as R
              where T.course_id= R.course_id
                 and R.year = 2009);
```

## Subqueries in the Form Clause

Find the average instructors' salaries of those departments where the average salary is greater than \$42,000."

```
select dept_name, avg_salary
from (select dept_name, avg (salary) as avg_salary
      from instructor
      group by dept_name) as xxx
where avg_salary > 42000;
```

Note that we do not need to use the **having** clause

## With Clause CTE

The **with** clause provides a way of defining a temporary relation whose definition is available only to the **query** in which the **with** clause occurs.

Find all departments with the maximum budget

```
with max_budget (value) as
  (select max(budget)
   from department)
select department.name
from department, max_budget
where department.budget = max_budget.value;
```

# Subqueries in the Select Clause

Scalar subquery is one which is used where a single value is expected

List all departments along with the number of instructors in each department

```
select dept_name,  
       (select count(*)  
        from instructor  
        where department.dept_name = instructor.dept_name)  
       as num_instructors  
from department;
```

Runtime error if subquery returns more than one result tuple

# Join operations – Example

Relation *course*

| <i>course_id</i> | <i>title</i> | <i>dept_name</i> | <i>credits</i> |
|------------------|--------------|------------------|----------------|
| BIO-301          | Genetics     | Biology          | 4              |
| CS-190           | Game Design  | Comp. Sci.       | 4              |
| CS-315           | Robotics     | Comp. Sci.       | 3              |

Relation *prereq*

| <i>course_id</i> | <i>prereq_id</i> |
|------------------|------------------|
| BIO-301          | BIO-101          |
| CS-190           | CS-101           |
| CS-347           | CS-101           |

Observe that

prereq information is missing for CS-315 and  
course information is missing for CS-437

# Left Outer Join

*course* **natural left outer join** *prereq*

| <i>course_id</i> | <i>title</i> | <i>dept_name</i> | <i>credits</i> |
|------------------|--------------|------------------|----------------|
| BIO-301          | Genetics     | Biology          | 4              |
| CS-190           | Game Design  | Comp. Sci.       | 4              |
| CS-315           | Robotics     | Comp. Sci.       | 3              |

| <i>course_id</i> | <i>prereq_id</i> |
|------------------|------------------|
| BIO-301          | BIO-101          |
| CS-190           | CS-101           |
| CS-347           | CS-101           |

| <i>course_id</i> | <i>title</i> | <i>dept_name</i> | <i>credits</i> | <i>prereq_id</i> |
|------------------|--------------|------------------|----------------|------------------|
| BIO-301          | Genetics     | Biology          | 4              | BIO-101          |
| CS-190           | Game Design  | Comp. Sci.       | 4              | CS-101           |
| CS-315           | Robotics     | Comp. Sci.       | 3              | <i>null</i>      |



# Right Outer Join

*course* **natural right outer join** *prereq*

| <i>course_id</i> | <i>title</i> | <i>dept_name</i> | <i>credits</i> |
|------------------|--------------|------------------|----------------|
| BIO-301          | Genetics     | Biology          | 4              |
| CS-190           | Game Design  | Comp. Sci.       | 4              |
| CS-315           | Robotics     | Comp. Sci.       | 3              |

| <i>course_id</i> | <i>prereq_id</i> |
|------------------|------------------|
| BIO-301          | BIO-101          |
| CS-190           | CS-101           |
| CS-347           | CS-101           |

| <i>course_id</i> | <i>title</i> | <i>dept_name</i> | <i>credits</i> | <i>prereq_id</i> |
|------------------|--------------|------------------|----------------|------------------|
| BIO-301          | Genetics     | Biology          | 4              | BIO-101          |
| CS-190           | Game Design  | Comp. Sci.       | 4              | CS-101           |
| CS-347           | <i>null</i>  | <i>null</i>      | <i>null</i>    | CS-101           |

# Outer Join

- An extension of the join operation that avoids loss of information.
- Computes the join and then adds tuples from one relation that does not match tuples in the other relation to the result of the join.
- Uses *null* values.

## Joined Relations

- **Join operations** take two relations and return as a result another relation.
- These additional operations are typically used as subquery expressions in the **from** clause
- **Join condition** – defines which tuples in the two relations match, and what attributes are present in the result of the join.
- **Join type** – defines how tuples in each relation that do not match any tuple in the other relation (based on the join condition) are treated.

| <i>Join types</i>   | <i>Join Conditions</i>   |
|---|--|
| inner join<br>left outer join<br>right outer join<br><b>full outer join</b> | <b>natural</b><br>on < <b>predicate</b> ><br><b>using</b> ( $A_1, A_1, \dots, A_n$ ) |

# Full Outer Join

*course* **natural full outer join** *prereq*

| <i>course_id</i> | <i>title</i> | <i>dept_name</i> | <i>credits</i> |
|------------------|--------------|------------------|----------------|
| BIO-301          | Genetics     | Biology          | 4              |
| CS-190           | Game Design  | Comp. Sci.       | 4              |
| CS-315           | Robotics     | Comp. Sci.       | 3              |

| <i>course_id</i> | <i>prereq_id</i> |
|------------------|------------------|
| BIO-301          | BIO-101          |
| CS-190           | CS-101           |
| CS-347           | CS-101           |

| <i>course_id</i> | <i>title</i> | <i>dept_name</i> | <i>credits</i> | <i>prereq_id</i> |
|------------------|--------------|------------------|----------------|------------------|
| BIO-301          | Genetics     | Biology          | 4              | BIO-101          |
| CS-190           | Game Design  | Comp. Sci.       | 4              | CS-101           |
| CS-315           | Robotics     | Comp. Sci.       | 3              | <i>null</i>      |
| CS-347           | <i>null</i>  | <i>null</i>      | <i>null</i>    | CS-101           |

# Credits

- The SQL slides come from:
- Silberschatz, Korth and Sudarshan, Database Systems Concepts 6th Edition, Ch3.