Query Planning:
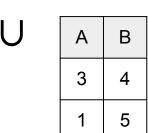## Translating SQL into Relational Algebra
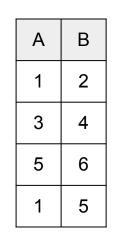
ULB

Refreshing the Relational Algebra

- Relations are tables whose columns have names, called attributes

- The set of all attributes of a relation is called the schema of the relation

- The rows in a relation are called tuples

- A relation is set-based if it does not contain duplicate tuples.

- It is called bag-based otherwise.

- A Relational Algebra (RA) operator takes as input 1 or more relations

  and produces as output a new relation

A	B	C	D
1	2	3	4
1	2	3	5
3	4	5	6
5	6	3	4

ULB

Union ∪ / Intersection ∩ / Difference -

Set-based   Bag-based

A	B
1	2
3	4
5	6

∪

A	B
3	4
1	5

=

A	B
1	2
3	4
5	6
1	5

A	B
1	2
3	4
5	6
1	5
3	4

- Input relations must have the same schema (same set of attributes)
- Historically speaking, relations are defined to be sets of tuples: duplicate tuples cannot occur in a relation.
- In practical systems, however, it is more efficient to allow duplicates to occur in relations, and only remove duplicates when requested. In this case relations are bags.

ULB

Selection

$$\sigma_{A>=3}$$

A	B
1	2
3	4
5	6

=

A	B
3	4
5	6

ULB

Projection

$$\pi_{A,C}$$

A	B	C	D
1	2	3	5
3	4	3	6
5	6	5	9
1	6	3	5

=

Set-based

A	C
1	3
3	3
5	5

ULB

Cartesian Product

A	B
1	2
3	4

✗

C	D
2	6
3	7
4	9

=

A	B	C	D
1	2	2	6
1	2	3	7
1	2	4	9
3	4	2	6
3	4	3	7
3	4	4	9

Input relations must have disjoint schema (disjoint set of attributes), otherwise rename first

ULB

Natural Join

Join on the commun attributes

A	B
1	2
3	4

⋈

B	D
2	6
3	7
4	9

=

A	B	D
1	2	6
3	4	9

A B C

B C D

A B C D X.B=Y.B X.C=Y.C

ULB

Natural Join

A	B
1	2
3	4

⋈

C	D
2	6
3	7
4	9

=

A	B	C	D
1	2	2	6
1	2	3	7
1	2	4	9
3	4	2	6
3	4	3	7
3	4	4	9

Same as cartesian product

ULB

Theta Join

A	B
1	2
3	4

$\bowtie_{B=C}$

C	D
2	6
3	7
4	9

=

A	B	C	D
1	2	2	6
3	4	4	9

ULB

Renaming

$$\rho_T \quad \begin{array}{|c|c|} \hline A & B \\ \hline 1 & 2 \\ \hline 3 & 4 \\ \hline 5 & 6 \\ \hline \end{array} \quad = \quad \begin{array}{|c|c|} \hline T.A & T.B \\ \hline 1 & 2 \\ \hline 3 & 4 \\ \hline 5 & 6 \\ \hline \end{array}$$

Renaming specifies that the input relation (and its attributes) should be given a new name.

Relational Algebra Expressions

Built using relation variable, AND

RA operators

$$\sigma_{\text{length}>=100}(\text{Movie}) \bowtie_{\text{title=movietitle}} \text{StarsIn}$$

Write the equivalent SQL

select * from Movie, StarsIn where title=movietitle and length>=100;

ULB

The Extended Relational Algebra

Add more operators

Extended projection

allows renaming

π

A,C->D

Set-based

A	B	C	D
1	2	3	5
3	4	3	6
5	6	5	9
1	6	3	5

=

A	D
1	3
3	3
5	5

ULB

The Extended Relational Algebra

Add more operators

Grouping

γ $_{A,min(B)->D}$

A	B	C
1	2	a
1	3	b
2	3	c
2	4	a
2	5	a

=

A	D
1	2
2	3

ULB

The Extended Relational Algebra

	Operator		Operator
∪	Union	⋈	Natural join
∩	Intersection	⋈$_{B=C}$	Theta join
-	Difference	⋈$_{B=C}$	Left outer join
σ$_{A>=3}$	Selection	⋈$_{B=C}$	Right outer join
π$_{A,C}$	Projection	⋈$_{B=C}$	Full outer join
×	Cartesian product	γ$_{A,min(B)->D}$	Aggregation
ρ	Rename	←	Assignment

Query Planning

Query Translation

Query Translation

SQL → [Lexical Analysis] → Stream of tokens → [Syntactic analysis] → Abstract Syntax Tree → [Transformation] → Logical query plan "Intermediate code"

ULB

Translating SQL into Relational Algebra

In the examples that follow, we will use the following database:

- Movie(title: string, year: int, length: int, genre: string, studioName: string,

  producerC#: int)

- MovieStar(name: string, address: string, gender: char, birthdate: date)

- StarsIn(movieTitle: string, movieYear: string, starName: string)

- MovieExec(name: string, address: string, CERT#: int, netWorth: int)

- Studio(name: string, address: string, presC#: int)

ULB

select-from-where

```
SQL:
SELECT movieTitle
FROM    StarsIn S, MovieStar M
WHERE   S.starName = M.name AND M.birthdate = 1960
```

RA ?

ULB

select-from-where

SQL:
SELECT movieTitle
FROM   StarsIn S, MovieStar M
WHERE   S.starName = M.name AND M.birthdate = 1960

RA ?

$\pi_{movieTitle} \ \sigma_{S.starName=M.name \ and \ M.birthdate=1960}(\rho_S(StarsIn) \times \rho_M(MovieStar))$

Other translations ?

ULB

select-from-where-groupby

```
SQL:
SELECT movieTitle, count(S.startName) AS numStars
FROM   StarsIn S, MovieStar M
WHERE  S.starName = M.name
GROUP BY movieTitle
```

RA ?

select-from-where-groupby

SQL:
SELECT movieTitle, count(S.startName) AS numStars
FROM   StarsIn S, MovieStar M
WHERE   S.starName = M.name
GROUP BY movieTitle

RA ?

$\gamma_{\text{M.movieTitle,count(S.starName)->numStars}}($

$\rho_{S}(\text{StarsIn})\bowtie_{\text{S.starName=M.name}}\rho_{M}(\text{MovieStar}))$

ULB

select-from-where-groupby-having

```
SELECT movieTitle, count(S.startName) AS numStars
FROM    StarsIn S, MovieStar M
WHERE   S.starName = M.name
GROUP BY movieTitle
HAVING count(S.startName) > 5
```

RA ?

ULB

select-from-where-groupby-having

```
SELECT movieTitle, count(S.startName) AS numStars
FROM    StarsIn S, MovieStar M
WHERE   S.starName = M.name
GROUP BY movieTitle
HAVING count(S.startName) > 5
```

RA ?

Relation

$$\sigma_{numStarts>5}(\gamma_{M.movieTitle,count(S.starName)->numStars}($$

$$\rho_S(StarsIn)\bowtie_{S.starName=M.name}\rho_M(MovieStar)))$$

ULB

Subqueries

```
SELECT *
FORM huge
WHERE c1 IN
    (SELECT c1 FROM tiny)
```

Loop over huge (optimizer has no choice)

V.S.

```
SELECT *
FORM huge h, tiny t
WHERE h.c1=t.c1
```

Optimizer has more choices

Which query is better ?

ULB

foreign
jit
lib
libpq
main
nodes
▼ optimizer
  ▶ geqo
  ▶ path
  ▼ plan
    ▶ analyzejoins.c
    ▶ createplan.c
    ▶ initsplan.c
    ▶ planagg.c
    ▶ planmain.c
    ▶ planner.c
    ▶ setrefs.c
    ▶ subselect.c
  ▶ prep
  ▶ util
▶ parser

```
645      /*
646       * If there is a WITH list, process each WITH query and either convert it
647       * to RTE_SUBQUERY RTE(s) or build an initplan SubPlan structure for it.
648       */
649      if (parse->cteList)
650          SS_process_ctes(root);
651
652      /*
653       * If the FROM clause is empty, replace it with a dummy RTE_RESULT RTE, so
654       * that we don't need so many special cases to deal with that situation.
655       */
656      replace_empty_jointree(parse);
657
658      /*
659       * Look for ANY and EXISTS SubLinks in WHERE and JOIN/ON clauses, and try
660       * to transform them into joins.  Note that this step does not descend
661       * into subqueries; if we pull up any subqueries below, their SubLinks are
662       * processed just before pulling them up.
663       */
664      if (parse->hasSubLinks)
665          pull_up_sublinks(root);          ⬅
666
667      /*
668       * Scan the rangetable for function RTEs, do const-simplification on them,
669       * and then inline them if possible (producing subqueries that might get
670       * pulled up next).  Recursion issues here are handled in the same way as
671       * for SubLinks.
672       */
673      preprocess_function_rtes(root);
674
675      /*
676       * Check to see if any subqueries in the jointree can be merged into this
677       * query.
678       */
679      pull_up_subqueries(root);          ⬅
680
```

ULB

Subquery processing and transformations

Subqueries are notoriously expensive to evaluate. This section describes some of the transformations that Derby makes internally to reduce the cost of evaluating them.

Materialization
Materialization means that a subquery is evaluated only once. There are several types of subqueries that can be materialized.

Flattening a subquery into a normal join

Flattening a subquery into an EXISTS join

Flattening VALUES subqueries

DISTINCT elimination in IN, ANY, and EXISTS subqueries

IN/ANY subquery transformation

Parent topic: Internal language transformations

Related concepts
Predicate transformations
Transitive closure

Subqueries

We can always normalize subqueries to use only EXISTS and NOT EXISTS [Van den Bussche, Vansummeren] [1,2]

```
SELECT movieTitle FROM StarsIn
WHERE starName IN (SELECT name
                       FROM MovieStar
                       WHERE birthdate=1960)


⇒     SELECT movieTitle FROM StarsIn
      WHERE EXISTS (SELECT name
                       FROM MovieStar
                       WHERE birthdate=1960 AND name=starName)
```

1 Only valid for set-based Relations
2 https://cs.ulb.ac.be/public/_media/teaching/infoh417/sql2alg_eng.pdf

ULB

Subqueries

We can always normalize subqueries to use only EXISTS and NOT EXISTS [Van den Bussche, Vansummeren] [1,2]

```
SELECT name FROM MovieExec
WHERE netWorth >= ALL (SELECT E.netWorth
                              FROM MovieExec E)
```

```
⇒    SELECT name FROM MovieExec
     WHERE NOT EXISTS(SELECT E.netWorth
                            FROM MovieExec E
                            WHERE netWorth < E.netWorth)
```

1 Only valid for set-based Relations
2 https://cs.ulb.ac.be/public/_media/teaching/infoh417/sql2alg_eng.pdf

ULB

Subqueries

We can always normalize subqueries to use only EXISTS and NOT EXISTS [Van den Bussche, Vansummeren] [1,2]

```
SELECT C FROM S
WHERE C IN (SELECT SUM(B) FROM R
                GROUP BY A)



    ⇒    ?
```

1 Only valid for set-based Relations
2 https://cs.ulb.ac.be/public/_media/teaching/infoh417/sql2alg_eng.pdf

ULB

Subqueries

We can always normalize subqueries to use only EXISTS and NOT EXISTS [Van den Bussche, Vansummeren] [1,2]

```
SELECT C FROM S
WHERE C IN (SELECT SUM(B) FROM R
                GROUP BY A)
```

```
⇒    SELECT C FROM S
     WHERE EXISTS (SELECT SUM(B) FROM R
                    GROUP BY A
                    HAVING SUM(B) = C)
```

1 Only valid for set-based Relations
2 https://cs.ulb.ac.be/public/_media/teaching/infoh417/sql2alg_eng.pdf

ULB

Normalization

- Before translating a query we first <mark>normalize</mark> it such that all of the subqueries that occur in a WHERE condition are of the form EXISTS or NOT EXISTS.

- We may hence assume without loss of generality in what follows that all subqueries

ULB

Correlated Subqueries

A subquery can refer to attributes of relations that are introduced in an outer query.

```
SELECT movieTitle
FROM StarsIn S
WHERE EXISTS (SELECT name
              FROM MovieStar
              WHERE birthdate=1960 AND name=S.starName)
```

- The "outer" relations are called the context relations of the subquery.

- The set of all attributes of all context relations of a subquery are called the parameters of the subquery.

ULB

Translation of correlated select-from-where subqueries

```
SELECT S.movieTitle, M.studioName
FROM StarsIn S, Movie M
WHERE S.movieYear >= 2000
AND S.movieTitle = M.title
AND EXISTS (SELECT name
            FROM MovieStar
            WHERE birthdate=1960 AND name= S.starName)
```

ULB

Translation of correlated select-from-where subqueries

```
SELECT S.movieTitle, M.studioName
FROM StarsIn S, Movie M
WHERE S.movieYear >= 2000
AND S.movieTitle = M.title
AND EXISTS (SELECT name
            FROM MovieStar
            WHERE birthdate=1960 AND name= S.starName)
```

First translate the subquery

$$\pi_{name} \ \sigma_{birthdate=1960 \wedge name=S.starName}(MovieStar)$$

ULB

Translation of correlated select-from-where subqueries

```
SELECT S.movieTitle, M.studioName
FROM StarsIn S, Movie M
WHERE S.movieYear >= 2000
AND S.movieTitle = M.title
AND EXISTS (SELECT name
            FROM MovieStar
            WHERE birthdate=1960 AND name= S.starName)
```

Fix: add the context relation and parameters

$$\pi_{name,S.movieTitle,S.movieYear,S.starName}$$
$$\sigma_{birthdate=1960 \wedge name=S.starName}(MovieStar \times \rho_S (StarsIn))$$

ULB

Translation of correlated select-from-where subqueries

```
SELECT S.movieTitle, M.studioName
FROM StarsIn S, Movie M
WHERE S.movieYear >= 2000
AND S.movieTitle = M.title
AND EXISTS (SELECT name
            FROM MovieStar
            WHERE birthdate=1960 AND name= S.starName)
```

Next, translate the FROM clause of the outer query

$$\rho_S(\text{StarsIn}) \times \rho_M(\text{Movie})$$

Translation of correlated select-from-where subqueries

```
SELECT S.movieTitle, M.studioName
FROM StarsIn S, Movie M
WHERE S.movieYear >= 2000
AND S.movieTitle = M.title
AND EXISTS (SELECT name
            FROM MovieStar
            WHERE birthdate=1960 AND name= S.starName)
```

Synchronize both expressions by means of a join.

$$\rho_S(StarsIn) \times \rho_M(Movie) \bowtie$$
$$\pi_{name,S.movieTitle,S.movieYear,S.starName}$$
$$\sigma_{birthdate=1960 \wedge name=S.starName}(MovieStar \times \rho_S(StarsIn))$$

ULB

Translation of correlated select-from-where subqueries

```
SELECT S.movieTitle, M.studioName
FROM StarsIn S, Movie M
WHERE S.movieYear >= 2000
AND S.movieTitle = M.title
AND EXISTS (SELECT name
                FROM MovieStar
                WHERE birthdate=1960 AND name= S.starName)
```

Simplify

$\rho_M(\text{Movie}) \bowtie$

$\quad \pi_{S.movieTitle,S.movieYear,S.starName}$

$\quad\quad \sigma_{birthdate=1960 \wedge name=S.starName}(\text{MovieStar} \times \rho_S(\text{StarsIn}))$

ULB

Translation of correlated select-from-where subqueries

```
SELECT S.movieTitle, M.studioName
FROM StarsIn S, Movie M
WHERE S.movieYear >= 2000
AND S.movieTitle = M.title
AND EXISTS (SELECT name
            FROM MovieStar
            WHERE birthdate=1960 AND name= S.starName)
```

Complete the expression

$\pi_{\text{S.movieTitle,M.studioName}}$ $\sigma_{\text{S.movieYear>=2000}\wedge\text{S.movieTitle=M.title}}$

$\rho_{M}(\text{Movie})\bowtie$

$\pi_{\text{S.movieTitle,S.movieYear,S.starName}}$

$\sigma_{\text{birthdate=1960}\wedge\text{name=S.starName}}(\text{MovieStar} \times$

Wait !

```
SELECT S.movieTitle, M.studioName
FROM StarsIn S, Movie M
WHERE S.movieYear >= 2000
AND S.movieTitle = M.title
AND EXISTS (SELECT name
            FROM MovieStar
            WHERE birthdate=1960 AND name= S.starName)
```

SQL Result ?

```
Movie
=====
title  studioName  movieYear
----------------------
DBSA      ULB        2005


StartsIn
========
starName   movieTitle
----------------------
Foo           DBSA


MovieStar
========
name firstname birthdate
-------------------------
Foo    Bar       1960
Foo    Baz       1960
```

ULB

Wait !

```
SELECT S.movieTitle, M.studioName
FROM StarsIn S, Movie M
WHERE S.movieYear >= 2000
AND S.movieTitle = M.title
AND EXISTS (SELECT name
              FROM MovieStar
              WHERE birthdate=1960 AND name= S.starName)
```

SQL Result ?

```
movieTitle studioName
---------------------------------
DBSA         ULB
```

```
Movie
=====
title studioName  movieYear
----------------------
DBSA      ULB         2005


StartsIn
========
starName  movieTitle
----------------------
Foo           DBSA


MovieStar
========
name firstname birthdate
-------------------------
Foo   Bar       1960
Foo   Baz       1960
```

ULB

Wait !

$\pi_{\text{S.movieTitle,M.studioName}}$
  $\sigma_{\text{S.movieYear>=2000}\wedge\text{S.movieTitle=M.title}}$
    $\rho_M(\text{Movie}) \bowtie$
      $\pi_{\text{S.movieTitle,S.movieYear,S.starName}}$
        $\sigma_{\text{birthdate=1960}\wedge\text{name=S.starName}}($
          $\text{MovieStar} \times \rho_S(\text{StarsIn}))$

RA Result ?

```
Movie
=====
title   studioName   movieYear
----------------------
DBSA        ULB         2005


StartsIn
========
starName    movieTitle
----------------------
Foo             DBSA


MovieStar
========
name  firstname  birthdate
-------------------------
Foo    Bar        1960
Foo    Baz        1960
```

ULB

Wait !

$\pi_{S.movieTitle,M.studioName}$

  $\sigma_{S.movieYear>=2000 \wedge S.movieTitle=M.title}$

    $\rho_M(Movie) \bowtie$

      $\pi_{S.movieTitle,S.movieYear,S.starName}$

        $\sigma_{birthdate=1960 \wedge name=S.starName}($

          $MovieStar \times \rho_S(StarsIn))$

RA Result ?

```
movieTitle studioName
-------------------------------
DBSA        ULB
DBSA        ULB
```

Only for set-based flatterning

```
Movie
=====
title studioName  movieYear
-----------------------
DBSA       ULB        2005


StarsIn
========
starName   movieTitle
-----------------------
Foo           DBSA


MovieStar
========
name firstname birthdate
------------------------
Foo   Bar       1960
Foo   Baz       1960
```

ULB

Wait !

$\pi_{\text{S.movieTitle,M.studioName}}$

$\quad\sigma_{\text{S.movieYear>=2000}\wedge\text{S.movieTitle=M.title}}$

$\quad\quad\rho_M(\text{Movie})\bowtie$

$\quad\quad\quad\pi_{\text{S.movieTitle,S.movieYear,S.starName}}$

$\quad\quad\quad\quad\sigma_{\text{birthdate=1960}\wedge\text{name=S.starName}}($

$\quad\quad\quad\quad\quad\text{MovieStar} \times \rho_S(\text{StarsIn}))$

RA Result ?

```
movieTitle   studioName
--------------------------------
DBSA         ULB
DBSA         ULB
```

Only set-based Relations

```
Movie
=====
title  studioName   movieYear
----------------------------
DBSA        ULB          2005


StarsIn
=======
starName    movieTitle
----------------------------
Foo             DBSA


MovieStar
=========
name firstname birthdate
--------------------------
Foo    Bar        1960
Foo    Baz        1960
```

ULB

Flattening Subqueries in Bag-based Relations (probably all vendor implementations)

The requirements for flattening into a normal join are:

- **There is a uniqueness condition that ensures that the subquery does not introduce any duplicates if it is flattened into the outer query block.**

- **Each table in the subquery's FROM list (after any view, derived table, or subquery flattening) must be a base table.**

- **The subquery is not under an OR.**

- **The subquery is not in the SELECT list of the outer query block.**

- **The subquery type is EXISTS, IN, or ANY, or it is an expression subquery on the right side of a comparison operator.**

https://docs.oracle.com/javadb/10.8.3.0/tuning/ctuntransform36368.html

ULB

Flattening Subqueries in Bag-based Relations (probably all vendor implementations)

- There are no aggregates in the SELECT list of the subquery.

- The subquery does not have a GROUP BY clause.

- The subquery does not have an ORDER BY, result offset, or fetch first clause.

- If there is a WHERE clause in the subquery, there is at least one table in the subquery whose columns are in equality predicates with expressions that do not include any column references from the subquery block. These columns must be a superset of the key columns for any unique index on the table. For all other tables in the subquery, the columns in equality predicates with expressions that do not include columns from the same table are a superset of the unique columns for any unique index on the table.

ULB

System R: Relational Approach to Database Management

M. M. ASTRAHAN, M. W. BLASGEN, D. D. CHAMBERLIN,
K. P. ESWARAN, J. N. GRAY, P. P. GRIFFITHS,
W. F. KING, R. A. LORIE, P. R. McJONES, J. W. MEHL,
G. R. PUTZOLU, I. L. TRAIGER, B. W. WADE, AND V. WATSON

IBM Research Laboratory

To read before the next lecture. We will discuss it in the lecture. Only read until end of The Optimizer section (unless you fall in love with it)

https://www.seas.upenn.edu/~zives/cis650/papers/System-R.PDF

ULB

Credits

Many slides are copied from:

- Stijn Vansummeren, Database Systems Architecture course slides.